

Models of Interception Movements

From Raw Trajectories to Individual Behavioural Signatures

Preliminary Results

Seman Libbiss & Paz Flashner · Workshop on Deep Learning

In collaboration with Prof. Jason Friedman, Tel Aviv University

1. Overview and goal

The goal of this project is to compress each interception movement into a small set of numbers — a low-dimensional latent code — that can (a) reconstruct the movement and (b) act as an individual “fingerprint” whose aggregate over a subject’s trials predicts their behaviour. We build up through three phases: a K-Means clustering check (is any subject structure present at all?), a polynomial-spline baseline (a non-machine-learning reference), and a Conditional Variational Autoencoder (CVAE), the core model. All numbers below are evaluated on **7 subjects held out entirely from training**, so every score measures generalisation to new people. This is a preliminary report (3-seed averages; a 10-seed run is in progress and will tighten the error bars without changing the conclusions).

2. Data and trial segmentation

We analyse the free eye-movement condition (condition 2): 3D finger trajectories at 240 Hz. We segment each trial by **task events** rather than by a velocity threshold. The target appears, holds still for a randomised foreperiod (0.18–0.48 s), then **starts moving** (the go-signal); the participant may only move after that. We take the **go-signal as the behavioural zero-time** (so reaction time is measured from it and the randomised foreperiod does not leak into the data), and the window **ends at finger arrival** (interception). Each movement is low-pass filtered (10 Hz), resampled to 100 frames, and origin-aligned.

From 4,763 condition-2 trials we retain **4,684 (98.3%)**, dropping 48 “too early” (moved before the go-signal), 27 that arrive >1 s after the target’s window (disengagement), and 4 timeouts. The event-based window also removed a segmentation artefact in which late sensor jitter had stretched a 0.5 s reach to as long as 9.7 s; the maximum movement time is now 1.36 s.

3. Phase 1 — K-Means baseline: is the signal even there?

K-Means is a classical clustering algorithm: it groups trials by geometric similarity, with no notion of “subject.” We check whether the clusters line up with the true subjects (Adjusted Rand Index, ARI: 1.0 = perfect match, 0 = random).

Representation	ARI	NMI
Raw trajectories (300-dim)	0.093	0.292
Kinematic features (11-dim)	0.104	0.294

Interpretation. Both scores are weak but clearly above chance (0). A little subject structure exists, but a simple algorithm barely recovers it. This is a *good* baseline result: it confirms individual signatures are real yet hard to extract, which is exactly what motivates a deep generative model. (A high K-Means score would undermine the need for one.)

4. Phase 2 — Polynomial spline baseline

The spline is pure mathematics (no learning): we fit a smooth cubic curve to each trajectory. It copies the *shape* almost perfectly — a per-trial fit reaches a reconstruction MSE of ~ 0.0002 , an interpolation ceiling — and even a capacity-matched version (Spline+PCA, reduced to the same number of dimensions) reconstructs geometry slightly better than the CVAE (0.20 vs 0.28 at $n=3$). **But a spline's coefficients describe a line, not a person.** On the quantities we care about (timing, curvature for unseen subjects, and generating realistic movement distributions), the CVAE beats it significantly. The spline is a better *tracer*; the CVAE is a better *portrait painter*. This is why reconstruction MSE alone is the wrong scoreboard for a fingerprinting project.

5. Phase 3 — the Conditional VAE (our model)

What it is and what it outputs

The CVAE has an encoder that squeezes a trajectory into n latent numbers, and a decoder that rebuilds it. We chose a **variational** autoencoder (not a plain one) because its latent space is smooth and continuous — which is what makes averaging a subject's trials into one meaningful fingerprint valid. It is **conditional**: we feed the task facts (start position, side) separately, so the latent is free to encode personal style rather than the task. The decoder has two outputs: **the full reconstructed trajectory** (its main job) **and a timing head** that predicts movement time and reaction time — the temporal axis that resampling removes (following Prof. Friedman's suggestion). So a generated code yields a shape *and* its durations.

Latent-dimension sweep

We swept the fingerprint size n over $\{2, 3, 4, 8, 16\}$ (3 seeds each).




Figure 1. As n grows, reconstruction improves, timing prediction saturates around the $n=8$ elbow ($R^2 \sim 0.99$), and the number of behavioural features captured rises.


n	Recon MSE	Timing R^2 (move)	Timing R^2 (react)	Behav. features > chance (of 11)
2	0.43	0.80	0.85	2.0
3	0.30	0.89	0.93	2.3
4	0.16	0.92	0.96	3.0
8	0.09	0.99	0.99	3.7
16	0.06	0.99	0.99	5.3

What we chose and why. We report all n , and take **$n = 3$ as the headline fingerprint**: it is the smallest size that is directly visualizable, already captures timing at $R^2 \sim 0.9$, and stays interpretable. **$n = 8$** is the accuracy elbow (reconstruction and timing saturate); **$n = 16$** captures the most behaviour but is not interpretable. **Compared with the spline**, the CVAE gives up a little point-wise geometry in exchange for

a structured, interpretable latent space and far stronger timing and behavioural prediction — the properties the fingerprint needs.


What each latent variable controls

Because $n=3$ is small, we can ask what each latent number does by correlating it with movement features (Figure 2). All three axes turn out to encode **timing and speed** (reaction time, movement time, peak speed) in different combinations; none clearly isolates spatial shape such as curvature. In other words, the fingerprint the model learns is dominated by *temporal* dynamics. This is useful and honest: it tells us the latent is interpretable but timing-centric, and points to spatial/strategy features (e.g. number of sub-movements) as something a future version could encode more explicitly.



6. Subject fingerprints

A subject **fingerprint** is the aggregate of that subject's trial latent codes — the mean position (who they are on average) plus the spread (their trial-to-trial variability). Probing from the fingerprint to a subject's behaviour, for unseen subjects, the number of behavioural features predicted above chance rises with capacity ($\sim 2/11$ at $n=3$ to $\sim 5/11$ at $n=16$), but per-feature accuracy stays weak and subjects do not separate cleanly.



This weakness is a data limitation, not a modelling failure. It persists across every fingerprint size, and an independent comparison found a 1,700 $\times$ -larger model and a small PCA model plateauing at the same separation. With only 28 subjects, and trial-to-trial variability larger than the differences between

people, the individual signature is only weakly recoverable. We therefore present it as a validated finding and lead with the strong kinematic and timing results.

7. The dashboard (deliverable)

We built an interactive tool on the trained model with two modes. **Exploration** lets a researcher move the latent sliders and watch a movement (and its timing) be generated — a way to see what the fingerprint space contains. **Inference** takes a subject's trials and reads out their fingerprint (mean + spread), reproducing Figure 3 for any uploaded subject.




Figure 4. Exploration mode: four latent settings produce four different generated movements, each with its own movement and wait time.

8. Where we followed the proposal, and where we deviated

For transparency, this compares what the submitted proposal said we would do against what we did.

Proposal item	Status	Note
K-Means baseline (2 representations)	Done	Raw trajectory + kinematic features.
Spline baseline (cubic, few knots)	Done	As specified.
CVAE + timing head	Done	Predicts movement time and reaction time.
Latent sweep $n = 2,3,4,8,16$	Done	Full sweep; headline $n=3$, elbow $n=8$.
Eval: MSE, Spearman, probing R^2 , KS, MMD, energy	Done	All implemented.
Dashboard (Inference + Exploration)	Done +	Extended to multi-trial subject fingerprints.
Segmentation of the movement window	Refined	Changed to event-based (go-signal to arrival), beyond the proposal, after investment.
Trial filtering	Added	Drop too-early / >1 s-late / timeout using the .mat outcome labels (not described in proposal).
Benchmark vs Friedman sub-movement pipeline	Not done	External dependency not yet available; pending (see questions).
Fallback to a supervised predictor	Not needed	The VAE converged; the fallback was not triggered.

9. Methods summary

- **Data:** condition 2, 4,684 trials, 28 subjects; 3D finger position at 240 Hz.
- **Preprocessing:** event-based segmentation (go-signal to arrival), 10 Hz Butterworth low-pass, resample to 100 frames (cubic spline), origin normalisation.
- **Split:** leave-N-subjects-out, 17 train / 4 validation / 7 test.
- **Model:** Conditional VAE, encoder/decoder MLP (hidden 256), condition = start position + side, timing head predicting movement & reaction time; KL annealing; standard-ELBO loss.
- **Baselines:** K-Means (trajectory & feature), polynomial spline (per-trial & Spline+PCA).
- **Evaluation (7 held-out subjects):** reconstruction MSE; timing R^2 ; latent-feature Spearman correlations; behavioural probing (linear + SVR, leave-one-subject-out); generative fidelity (KS, MMD, energy distance).
- **Tools:** Python, PyTorch, scikit-learn, NumPy/Pandas, Matplotlib, Streamlit.

10. Optional next steps (pending your feedback)

We are holding these until we hear back on the methods above, and note them as options rather than commitments:

- **Sequence model:** replace the MLP encoder/decoder with a 1-D convolutional (or recurrent) network that respects time-order. Likely to improve movement-shape reconstruction; unlikely to overcome the fingerprint data limit.
- **Exact target speed:** feed the precise per-trial target speed (recoverable from the stimulus files) as a condition, instead of only the speed range.
- **Sub-movement / peak-count feature:** explicitly encode or predict the number of sub-movements, to capture the hesitation strategy that the current timing-centric latent under-represents.
- **More subjects** would be the only thing likely to materially improve fingerprint separation.

11. Questions for Prof. Friedman

- **The 'not fixating' flag.** It appears on ~29% of condition-2 (free-eye) trials — more than in condition 1 — and those trials' hand movements look normal. Was there an eye tracker, and is that check meaningful in condition 2? We are currently keeping these trials.
- **Segmentation zero-time.** Do you agree the correct reference is the target's motion onset (the go-signal), and that predicting movement + reaction time addresses the resampling concern?
- **Late-arrival cutoff.** We drop trials arriving >1 s after the target's window as disengagement (there is no clean gap in the data). Is 1 s reasonable?
- **'Too early' trials** (moved before the go-signal) — we exclude these as invalid; please confirm.
- **Sub-movement benchmark.** Could we get access to your sub-movement decomposition pipeline to benchmark generative fidelity against it?